

CS2850 operating system lab
week 2: types, variables, functions

nicolo colombo
nicolo.colombo@rhul.ac.uk
Office Bedford 2-21

outline

variables

functions

types

composite objects

programs

C programs consist of **functions** and **variables**

a function is a set of **statements** that specify the operations to be performed

a variable is a **location in storage** with three attributes:

- an **identifier** or *name*
- a **storage class** determining the *lifetime* of the identified storage
- a **type** determining the *meaning* of the value found in the identified storage

declaration

`int main()` is a *special* function because the program starts executing from its first line

try to compile with `gcc -Wall -Werror -Wpedantic` a program with *more than one* or *without* any `main` to see what happens

in C functions and variables should always be **declared** before they are used (but you can separate declaration and definition)

example

```
#include <stdio.h>
void printInt(int j);
int main() {
    int i;
    i = 1;
    printInt(i);
}
void printInt(int k) {
    printf("i=%d\n", k);
}
```

1
2
3
4
5
6
7
8
9
10

note that, in C, the declaration of a variable **does not initialise** it automatically

comment out Line 5 to see what happens

variables

a variable is a *named storage location* with a specified *lifetime* and *meaning*

the variable properties are fixed in its *declaration*

```
storage_class variable_type variable_name;
```

the *lifetime* of a variable is specified by *where* it is declared and by its *storage class*

types and names

the **meaning** of a variable is determined by its *type* and specifies how the stored bytes should be *interpreted*

this is why the type of a variable should be *always* be specified

names (or *identifiers*) are sequences of letters and digits used in the program for referring to objects, i.e. names like `i` or `printInt` in the example above

storage classes

two important storage classes are

- **automatic**: the variable is local to the block and reinitialised every time the function is called
- **static**: the variable keeps its value across different function calls

by default, *local* variables are automatic and *global* variables are static

uninitialised automatic variables have an *undefined value*

you can add the *qualifier* `const` to the declaration of any variable to specify that its value **will not change**

Example

```
#include <stdio.h>
static int j = 0;
void printInt(int i) {
    i++;
    printf("i=%d\n", i);
    j++;
    printf("j=%d\n", j);
}
int main() {
    int i = 1;
    printInt(i);
    printInt(i);
}
```

try to predict the output of the program above and what happens if you add the qualifier `const` to the two variables?*

*always compile your programs with `gcc -Wall -Werror -Wpedantic`

functions

a function is **defined** by writing

```
return_type  function_name(arguments) {statements};
```

where

- `return_type` is the *type* of the value that the function returns
- `function_name` is the name of the function used to *call* it from other functions
- `arguments` is the name and type of the *parameters* that a calling function passes to the called function
- `statements` is the set of *operations* to be performed

function calls

a function is **called** by writing its name

```
f(argument_value_1, argument_value_2, \dots );
```

where `argument_value_I` is the value of parameter I (at the right place)

parameter names are not important in the function declaration and only their *type* is required (see the code above for an example)

example

this is the first code in Kernighan & Ritchie's [The C Programming Language](#)

```
    ]  
#include <stdio.h>  
main() {  
    printf("hello, world\n");  
}
```

if you compile the program with `gcc -Wall -Werror -Wpedantic` you get the *error* message

```
gcc -Wall -Werror -Wpedantic kernighan.c  
kernighan.c:3:1: error: return type defaults to 'int' [-Werror=implicit-int]  
    3  main() ^~~~  
cc1: all warnings being treated as errors
```

did Kernighan and Ritchie make a mistake?

notes

what is missing is the **return type** of the `main` function but the book was written before ANSI C conventions changed

`printf` can be **just called** because it is declared and defined in `stdio.h`

the **argument** list in `main` is empty

the argument **value** in `printf("hello, world\n")` is the string constant `"hello, world\n"`

types of C (1)

the type of a variable specifies the *meaning of the bytes* stored at the associated location

there are a few **basic** data types in C

- **char** a single character, e.g. a
- **int** an integer, e.g. -1234
- **unsigned int** a positive integer, e.g. 12345
- **float** a single-precision floating point, e.g. 23,923

for *more* types available in C, check the full list in appendix A4 of **The C Programming Language**

void refers to an empty set of values but it is not a regular type, e.g. you cannot declare a void variable by writing `void n;`

types of C (2)

the **number of bytes** associated with each type may be architecture/compiler specific

you can use the operator **sizeof()** to get the number of bytes of any given basic type

run this program to check the size of the basic C types on your system

```
#include <stdio.h>
int main() {
    printf("sizeof(char)=%lu\n", sizeof(char));
    printf("sizeof(int)=%lu\n", sizeof(int));
    printf("sizeof(unsigned int)=%lu\n", sizeof(unsigned int));
    printf("sizeof(float)=%lu\n", sizeof(float));
}
```

more on int and char

integer constants can be used to initialize a variable and given in binary (e.g. `int a = 0b01010101;`), decimal (`int a = 85;`), hexadecimal (`int a = 0x55;`) or octal (`int a = 0125;`)

character constants are characters enclosed in single quotes, (e.g. `char a = 'x';` or `char a = '\n'`)

the numerical value of a **char** is the **ASCII** integer code corresponding to it, e.g. check the output of

```
printf("ASCII(r)=%d\n", 'r');
```

printable characters, e.g. `r`, `$`, or `3`, are always positive but plain `char`'s can be **signed or unsigned** (depending on the machine)

conversions

the conversion of a 'true' character to a positive integer is called **integral promotion** in Appendix A6.1 of **The C Programming Language**

the conversion of an integer to an **unsigned int** (by left-truncation of its two's complement representation) is called **integral conversion** in Appendix A6.2

floating points are converted to integral types by discarding their fractional part

example

```
#include <stdio.h>
int main() {
    char a = -123;
    printf("character=%c\n", (char) a);
    printf("signed int=%d\n", (int) a);
    printf("signed int(octal)=%o\n", (int) a);
    printf("unsigend int=%u\n", (unsigned int) a);
    printf("unsigend int(octal)=%o\n", (unsigned int) a);
}
```

type casting specifiers as (unsigned int) force a *specified interpretation* of the following variable

the **output** of this program shows that -123 does not correspond to a *valid character* and that the *stored bytes* (printed in octal) can be interpreted in different ways

operators in C

arithmetic operators: +, -, *, /, %

relational operators: ==, !=, >, >=, <, <=

logical operators: && (and), || (or)

the **negation operator**, “!”, converts a non-zero operand into a 0, and a zero operand into 1 i.e. `if (!operand)` and `if(operand == 0)` are *equivalent*

% is the **modulus operator**, e.g. $1\%2 = 1$, $2\%1 = 0$, $3\%2 = 1$, and $2\%3 = 2$

bitwise operators (1)

C has an *"and"*, an *"or"*, two *"shift"*, and a *"complement"*
bitwise operators: `&`, `|`, `>>`, `<<`, `~`

bitwise operators can be used to manipulate *directly* the **binary representation** of variables

compile and run the program in the next slide to understand how they work

bitwise operators (2)

```
#include <stdio.h>
int main() {
    int i = 077;
    printf("i=%o\n", i);
    printf("~i=%o\n", ~i);
    printf("(i&i)=%o\n", i & i);
    printf("(i&~i)=%o\n", i & ~i);
    printf("(i|i)=%o\n", i | i);
    printf("(i|~i)=%o\n", i | ~i);
    printf("i<<1=%o\n", i<<1);
    printf("i<<8=%o\n", i<<8);
    printf("i>>1=%o\n", i>>1);
    printf("i>>8=%o\n", i>>8);
}
```

arithmetics

depending on the operands, operators may cause **conversion** of the value of an operand from one type to another

for example, **integer division** may or may not **truncate** any fractional part, e.g. $3/2 = 1$ but $3/2. = 1.5$ and $3./2 = 1.5$

have a look at the full list of **arithmetic conversion rules** in Appendix A6.5 of [The C Programming Language](#)

there are no explicit or implicit conversion rules between a **void** and non-void types

non-basic types

besides the basic types, there is a *conceptually infinite* class of **derived types**, which are built from the basic types in various ways

the most important *derived types* of C are

- **strings**: null-terminated lists of `char`'s,
- **arrays**: lists of objects of a given type,
- **functions**: sets of statements returning objects of a given type
- **pointers**: memory addresses of objects of a given type
- **structures**: general composite objects containing objects of different types

a few notes on composite objects

composition of objects can be run **recursively** e.g. you can have an *array of arrays*, include an array in a struct, or let a pointer point to an *array of pointers*

in C, the **structure** of composite objects, e.g. the number of entries of an array, should be declared before they are used and *cannot be changed at run time*

the **size** in bytes of composite objects can be obtained using **sizeof**

```
#include <stdio.h>
int main() {
    printf("sizeof(char[10])=%lu\n", sizeof(char[10]));
    printf("10 * sizeof(char)=%lu\n", 10 * sizeof(char));
}
```


strings (1)

the program in the next slide

- i) *declares and initialises* a string (constant)
- ii) *prints* the string using `printf` and the format specifier `%s`
- iii) prints a specific *character* of the string and non-initialised value *outside* the string
- iv) prints the string *starting at position 7*

you will see more about strings and understand all details of the program in the next weeks

strings (2)

```
#include <stdio.h>
int main() {
    char *s = "hello, world\n";
    printf("s=%s", s);
    printf("s[7]=%c\n", s[7]);
    printf("s[100]=%c\n", s[100]);
    printf("&s[7]=%s", &s[7]);
    //s[7] = 'x'; //uncomment this line if you declare s as a char array
}
```

1
2
3
4
5
6
7
8
9

the *ampersand* operator, & returns the address of a variable and here is used to *read the string starting at s[7]*

what happens if you replace the statement in Line 3 with

```
char s[] = "hello, world"
```

strings (3)

the program *does not crash* if you try to access **uninitialised entries** of `s` but the value stored there is *unpredictable*

you *cannot* change a specific entry of `s` because `s` is a string (and not an array of `char`'s)

the **format specifier** `%s` allows you to use `printf` for printing `s` with as a whole

the program *knows where the string terminates* because strings are **null-terminated** lists of `char`'s, i.e. the last `char` of a string is *always* a `'\0'`

arrays (1)

the program in the next slide

i) *declares* an array of 10 `int`, i.e. allocate the memory space to store 10 `int`

ii) *loads* random integers to its entries, through component-wise assignments

iii) prints the vector components with `printf` and the format specifier `%d`

C does not have built-in functions to manipulate arrays *as a unit*, e.g. you cannot print an array with a single call of `printf`

arrays are handled by referring to *the address or their first entry*

arrays (2)

```
#include <stdio.h>
#include <stdlib.h>
void loadVector(int a[], int size) {
    for (int i=0; i<size; i++) a[i] = ((float) 10 * rand())/RAND_MAX;
}
void printVector(int a[], int size) {
    for (int i=0; i<size; i++) printf("a[%d]=%d\n", i, a[i]);
}
int main() {
    int size = 10;
    int a[size];
    loadVector(a, size);
    printVector(a, size);
    printf("a[100]=%d\n", a[100]);
    a[7] = 100;
}
```

note that the changes made in `a` by `loadVector` are *not discarded* when the function returns

arrays (3)

the program *does not crash* if you try to access uninitialised entries of `a` but the value stored there is *unpredictable*

the program *does not block* the access to **unallocated** memory regions, i.e. `a[100]`

you can change a specific entry of `a` by adding a statement like `a[7] = '100';`

you need to write a customised function to load and print `a` as a single unit

`loadVector` and `printVector` do not know anything about the length of the input because `a` is passed as **the address of `a[0]`**